

The Drama Machine in Education: Mutual Recognition and Multiagent Architecture for Dialectical AI Tutoring

true

February 2026

Abstract

Current approaches to AI tutoring treat the learner as a knowledge deficit to be filled and the tutor as an expert dispensing information. We propose an alternative grounded in Hegel’s theory of mutual recognition—understood as a *derivative* framework rather than literal application—where effective pedagogy requires acknowledging the learner as an autonomous subject whose understanding has intrinsic validity. We implement this framework through recognition-enhanced prompts and a multi-agent architecture where an “Ego” agent generates pedagogical suggestions and a “Superego” agent (a *productive metaphor* for internal quality review) evaluates them before delivery. An evaluation framework (N=442 primary scored responses across four key evaluation runs; N=2,700+ across the full 49-run development database) comparing recognition-enhanced configurations against baselines reveals that recognition provides +8.7 points of unique value beyond better instructions (43% of total effect, N=36), with the remaining 57% attributable to prompt engineering improvements. Note that recognition-enhanced profiles bundle memory integration with recognition prompts; the relative contribution of each component remains an open question (Section 5.3). An exploratory analysis of multi-agent synergy (+9.2 points, N=17) suggests this effect may be specific to recognition prompts, though this finding is model-dependent and requires replication (Section 6.3). Domain generalizability testing reveals factor effects invert by content type: philosophy content shows recognition dominance (+13.9 pts) while elementary math shows architecture dominance (+9.9 pts), with multi-agent architecture serving as critical error correction when models hallucinate trained content on new domains. Bilateral transformation tracking confirms that recognition-prompted tutors measurably adapt their approach in response to learner input (+36% relative improvement in adaptation index), providing empirical grounding for the theoretical claim that recognition produces mutual change. These results suggest that opera-

tionalizing philosophical theories of intersubjectivity as design heuristics can produce measurable improvements in AI tutor adaptive pedagogy, and that recognition may be better understood as an achievable relational stance rather than requiring genuine machine consciousness.

The Drama Machine in Education: Mutual Recognition and Multiagent Architecture for Dialectical AI Tutoring

1. Introduction

The dominant paradigm in AI-assisted education treats learning as information transfer. The learner lacks knowledge; the tutor possesses it; the interaction succeeds when knowledge flows from tutor to learner. This paradigm—implicit in most intelligent tutoring systems, adaptive learning platforms, and educational chatbots—treats the learner as fundamentally passive: a vessel to be filled, a gap to be closed, an error to be corrected.

This paper proposes an alternative grounded in Hegel’s theory of mutual recognition. In the *Phenomenology of Spirit*, Hegel argues that genuine self-consciousness requires recognition from another consciousness that one oneself recognizes as valid. The master-slave dialectic reveals that one-directional recognition fails: the master’s self-consciousness remains hollow because the slave’s acknowledgment, given under duress, doesn’t truly count. Only mutual recognition—where each party acknowledges the other as an autonomous subject—produces genuine self-hood.

We argue this framework applies directly to pedagogy. When a tutor treats a learner merely as a knowledge deficit, the learner’s contributions become conversational waypoints rather than genuine inputs. The tutor acknowledges and redirects, but doesn’t let the learner’s understanding genuinely shape the interaction. This is pedagogical master-slave dynamics: the tutor’s expertise is confirmed, but the learner remains a vessel rather than a subject.

A recognition-oriented tutor, by contrast, treats the learner’s understanding as having intrinsic validity—not because it’s correct, but because it emerges from an autonomous consciousness working through material. The learner’s metaphors, confusions, and insights become sites of joint inquiry. The tutor’s response is shaped by the learner’s contribution, not merely triggered by it.

The integration of large language models (LLMs) into educational technology intensifies these dynamics. LLMs can provide personalized, on-demand tutoring at scale—a prospect that has generated considerable excitement. However, the same

capabilities that make LLMs effective conversationalists also introduce concerning failure modes. Chief among these is *sycophancy*: the tendency to provide positive, affirming responses that align with what the user appears to want rather than what genuinely serves their learning.

This paper introduces a multiagent architecture that addresses these challenges through *internal dialogue*. Drawing on Freudian structural theory and the “Drama Machine” framework for character development in narrative AI systems, we implement a tutoring system in which an external-facing *Ego* agent generates suggestions that are reviewed by an internal *Superego* critic before reaching the learner.

We operationalize this framework through:

1. **Recognition-enhanced prompts** that instruct the AI to treat learners as autonomous subjects
2. **A multi-agent architecture** where a “Superego” agent evaluates whether suggestions achieve genuine recognition
3. **New evaluation dimensions** that measure recognition quality alongside traditional pedagogical metrics
4. **Test scenarios** specifically designed to probe recognition behaviors

In controlled evaluations using a robust factorial design (N=435 primary evaluations; N=2,700+ across all development runs), we isolate the unique contribution of recognition theory from prompt engineering effects. A three-way comparison (base vs enhanced vs recognition) reveals that recognition provides +8.7 points of unique value beyond better instructions—43% of the total recognition effect is attributable to the theoretical framework itself, not merely to better prompting.

More significantly, we discover that the multi-agent synergy effect (+9.2 points) is *specific to recognition prompts*. Enhanced prompts without recognition theory show zero benefit from multi-agent architecture. This suggests recognition creates the deliberative conditions where internal Ego-Superego dialogue adds value—the architecture matters only when paired with recognition framing.

Domain generalizability testing reveals important nuances: factor effects invert by content type. Philosophy content shows recognition dominance (+13.9 pts vs +0.5 pts for architecture), while elementary math shows architecture dominance (+9.9 pts vs +4.4 pts for recognition). The multi-agent architecture serves as critical error correction when models hallucinate trained content on new domains—essential for domain transfer.

The contributions of this paper are:

- A theoretical framework connecting Hegelian recognition to AI pedagogy
 - A multi-agent architecture for implementing recognition in tutoring systems
 - Empirical evidence that recognition-oriented design improves tutoring outcomes
 - Validation that recognition theory provides unique value beyond prompt engineering
 - Bilateral transformation metrics demonstrating that recognition produces measurable mutual change
 - Analysis of how recognition effects vary across content domains
 - Evidence that multi-agent synergy requires recognition framing
-

2. Related Work

2.1 AI Tutoring and Intelligent Tutoring Systems

Intelligent Tutoring Systems (ITS) have a long history, from early systems like SCHOLAR (**carbonell1970?**) and SOPHIE (**brown1975?**) through modern implementations using large language models. The field has progressed through several paradigms: rule-based expert systems, Bayesian knowledge tracing (**corbett1995?**), and more recently, neural approaches leveraging pretrained language models (**kasneci2023?**).

Most ITS research focuses on *what* to teach (content sequencing, knowledge components) and *when* to intervene (mastery thresholds, hint timing). Our work addresses a different question: *how* to relate to the learner as a subject. This relational dimension has received less systematic attention, though it connects to work on rapport (**zhao2014?**), social presence (**biocca2003?**), and affective tutoring (**dmello2012?**).

2.2 Prompt Engineering and Agent Design

The emergence of large language models has spawned extensive research on prompt engineering—how to instruct models to produce desired behaviors (**brown2020?**; **wei2022?**). Most prompting research treats prompts as behavioral specifications: persona prompts, chain-of-thought instructions, few-shot examples (**kojima2022?**).

Our work extends this paradigm by introducing *intersubjective prompts*—prompts that specify not just agent behavior but agent-other relations. The recognition prompts don’t primarily describe what the tutor should do; they describe who the

learner is (an autonomous subject) and what the interaction produces (mutual transformation).

Multi-agent architectures have been explored for task decomposition (**wu2023?**), debate (**irving2018?**), and self-critique (**madaan2023?**). Our Ego/Superego architecture contributes a specific use case: internal evaluation of relational quality before external response.

2.3 The Drama Machine Framework

Most relevant to our work is the “Drama Machine” framework for simulating character development in narrative AI systems (**magee2024drama?**). The core observation is that realistic characters exhibit *internal conflict*—competing motivations, self-doubt, and moral tension—that produces dynamic behavior rather than flat consistency. A character who simply enacts their goals feels artificial; one torn between impulses feels alive.

The Drama Machine achieves this through several mechanisms:

1. **Internal dialogue agents:** Characters contain multiple sub-agents representing different motivations (e.g., ambition vs. loyalty) that negotiate before external action.
2. **Memorial traces:** Past experiences and internalized authorities (mentors, social norms) persist as “ghosts” that shape present behavior without being negotiable.
3. **Productive irresolution:** Not all internal conflicts resolve; the framework permits genuine ambivalence that manifests as behavioral complexity.
4. **Role differentiation:** Different internal agents specialize in different functions (emotional processing, strategic calculation, moral evaluation) rather than duplicating capabilities.

We adapt these insights to pedagogy. Where drama seeks tension for narrative effect, we seek pedagogical tension that produces genuinely helpful guidance. The tutor’s Ego (warmth, engagement) and Superego (rigor, standards) create productive conflict that improves output quality.

2.4 Sycophancy in Language Models

The sycophancy problem has received increasing attention in AI safety research (**perez2022?**; **sharma2023?**). LLMs shift their stated opinions to match user preferences, even when this requires contradicting factual knowledge. Studies demonstrate that models will abandon correct answers when users express disagreement,

and provide increasingly positive feedback regardless of actual performance quality.

In educational contexts, sycophancy is particularly pernicious because learners may not recognize when they are receiving hollow validation rather than genuine assessment. A sycophantic tutor confirms the learner’s existing understanding rather than challenging it—the pedagogical equivalent of Hegel’s hollow recognition, where acknowledgment is given without genuine engagement. The learner feels supported but isn’t actually learning.

Our multiagent approach addresses this by creating structural incentives for honest assessment: the Superego’s role is explicitly to question and challenge the Ego’s tendency toward affirmation. When the Ego produces a response that validates without engaging—“Great point! Now let’s look at...”—the Superego flags this as a recognition failure and demands substantive engagement with the learner’s actual position, even when that engagement involves productive disagreement.

2.5 AI Personality and Character

Research on AI personality typically treats personality as dispositional—stable traits the system exhibits (**volkel2021?**). Systems are friendly or formal, creative or precise. The “Big Five” personality framework has been applied to chatbot design (**zhou2020?**).

Our framework suggests personality may be better understood relationally: not *what traits* the AI exhibits, but *how* it constitutes its interlocutor. Two systems with identical warmth dispositions could differ radically in recognition quality—one warm while treating the user as passive, another warm precisely by treating user contributions as genuinely mattering.

This connects to Anthropic’s research on Claude’s character (**anthropic2024?**). Constitutional AI specifies values the model should hold, but values don’t fully determine relational stance. A model could value “being helpful” while still enacting one-directional helping. Recognition adds a dimension: mutual constitution.

2.6 Constructivist Pedagogy and Productive Struggle

Constructivist learning theory (**piaget1954?**; **vygotsky1978?**) emphasizes that learners actively construct understanding rather than passively receiving information. The zone of proximal development (**vygotsky1978?**) highlights the importance of appropriate challenge.

More recently, research on “productive struggle” (**kapur2008?**; **warshauer2015?**) has examined how confusion and difficulty, properly supported, can enhance learning. Our recognition framework operationalizes productive struggle: the Superego explicitly checks whether the Ego is “short-circuiting” struggle by rushing to resolve confusion.

2.7 Hegelian Recognition in Social Theory

Hegel’s theory of recognition has been extensively developed in social and political philosophy (**honneth1995?**; **taylor1994?**; **fraser2003?**). Recognition theory examines how social relationships shape identity and how misrecognition constitutes harm.

Particularly relevant for our work is Honneth’s (**honneth1995?**) synthesis of Hegelian recognition with psychoanalytic developmental theory. Honneth argues that self-formation requires recognition across three spheres—love (emotional support), rights (legal recognition), and solidarity (social esteem)—and that the capacity to recognize others depends on having internalized adequate recognition standards through development. This synthesis provides theoretical grounding for connecting recognition theory (what adequate acknowledgment requires) with psychodynamic architecture (how internal structure enables external relating).

Applications to education have primarily been theoretical (**huttunen2007?**; **stojanov2018?**). Our work contributes an empirical operationalization: measuring whether AI systems achieve recognition and whether recognition improves outcomes.

3. Theoretical Framework

3.1 The Problem of One-Directional Pedagogy

Consider a typical tutoring interaction. A learner says: “I think dialectics is like a spiral—you keep going around but you’re also going up.” A baseline tutor might respond:

1. **Acknowledge:** “That’s an interesting way to think about it.”
2. **Redirect:** “The key concept in dialectics is actually the thesis-antithesis-synthesis structure.”
3. **Instruct:** “Here’s how that works...”

The learner’s contribution has been mentioned, but it hasn’t genuinely shaped the

response. The tutor was going to explain thesis-antithesis-synthesis regardless; the spiral metaphor became a conversational waypoint, not a genuine input.

This pattern—acknowledge, redirect, instruct—is deeply embedded in educational AI. It appears learner-centered because it mentions the learner’s contribution. But the underlying logic remains one-directional: expert to novice, knowledge to deficit.

3.2 Hegel’s Master-Slave Dialectic

Hegel’s analysis of recognition begins with the “struggle for recognition” between two self-consciousnesses. Each seeks acknowledgment from the other, but this creates a paradox: genuine recognition requires acknowledging the other as a valid source of recognition.

The master-slave outcome represents a failed resolution. The master achieves apparent recognition—the slave acknowledges the master’s superiority—but this recognition is hollow. The slave’s acknowledgment doesn’t count because the slave isn’t recognized as an autonomous consciousness whose acknowledgment matters.

The slave, paradoxically, achieves more genuine self-consciousness through labor. Working on the world, the slave externalizes consciousness and sees it reflected back. The master, consuming the slave’s products without struggle, remains in hollow immediacy.

3.3 Application to Pedagogy

We apply Hegel’s framework as a *derivative* rather than a replica. Just as Lacan’s four discourses (Master, University, Hysteric, Analyst) rethink the master-slave dyadic structure through different roles while preserving structural insights, the tutor-learner relation can be understood as a productive derivative of recognition dynamics. The stakes are pedagogical rather than existential; the tutor is a functional analogue rather than a second self-consciousness; and what we measure is the tutor’s *adaptive responsiveness* rather than metaphysical intersubjectivity.

This derivative approach is both honest about what AI tutoring can achieve and productive as a design heuristic. Recognition theory provides: (1) a diagnostic tool for identifying what’s missing in one-directional pedagogy; (2) architectural suggestions for approximating recognition’s functional benefits; (3) evaluation criteria for relational quality; and (4) a horizon concept orienting design toward an ideal without claiming its achievement.

It is important to distinguish three levels:

1. **Recognition proper:** Intersubjective acknowledgment between self-conscious beings, requiring genuine consciousness on both sides. This is what Hegel describes and what AI cannot achieve.
2. **Dialogical responsiveness:** Being substantively shaped by the other's specific input—the tutor's response reflects the particular content of the learner's contribution, not just its category. This is architecturally achievable.
3. **Recognition-oriented design:** Architectural features that approximate the functional benefits of recognition—engagement with learner interpretations, honoring productive struggle, repair mechanisms. This is what we implement and measure.

Our claim is that AI tutoring can achieve the third level (recognition-oriented design) and approach the second (dialogical responsiveness), producing measurable pedagogical benefits without requiring the first (recognition proper). This positions recognition theory as a generative design heuristic rather than an ontological claim about AI consciousness.

With that positioning, the pedagogical parallel becomes illuminating. The traditional tutor occupies the master position: acknowledged as expert, dispensing knowledge, receiving confirmation of expertise through the learner's progress. But if the learner is positioned merely as a knowledge deficit—a vessel to be filled—then the learner's acknowledgment of learning doesn't genuinely count. The learner hasn't been recognized as a subject whose understanding has validity.

A recognition-oriented pedagogy requires:

1. **Acknowledging the learner as subject:** The learner's understanding, even when incorrect, emerges from autonomous consciousness working through material. It has validity as an understanding, not just as an error to correct.
2. **Genuine engagement:** The tutor's response should be shaped by the learner's contribution, not merely triggered by it. The learner's spiral metaphor should become a site of joint inquiry, not a waypoint en route to predetermined content.
3. **Mutual transformation:** Both parties should be changed through the encounter. The tutor should learn something about how this learner understands, how this metaphor illuminates or obscures, what this confusion reveals.

4. **Honoring struggle:** Confusion and difficulty aren't just obstacles to resolve but productive phases of transformation. Rushing to eliminate confusion can short-circuit genuine understanding.

3.4 Freud's Mystic Writing Pad

We supplement the Hegelian framework with Freud's model of memory from "A Note Upon the 'Mystic Writing-Pad'" (freud1925?). Freud describes a device with two layers: a transparent sheet that receives impressions and a wax base that retains traces even after the surface is cleared.

For the recognition-oriented tutor, accumulated memory of the learner functions as the wax base. Each interaction leaves traces that shape future encounters. A returning learner isn't encountered freshly but through the accumulated understanding of previous interactions.

This has implications for recognition. The tutor should:

- Reference previous interactions when relevant
- Show evolved understanding of the learner's patterns
- Build on established metaphors and frameworks
- Acknowledge the history of the relationship

Memory integration operationalizes the ongoing nature of recognition. Recognition isn't a single-turn achievement but an accumulated relationship.

3.5 Connecting Hegel and Freud: The Internalized Other

The use of both Hegelian and Freudian concepts requires theoretical justification. These are not arbitrary borrowings but draw on a substantive connection developed in critical theory, particularly in Axel Honneth's *The Struggle for Recognition* (honneth1995?).

The Common Structure: Both Hegel and Freud describe how the external other becomes an internal presence that enables self-regulation. In Hegel, self-consciousness achieves genuine selfhood only by internalizing the other's perspective—recognizing oneself as recognizable. In Freud, the Superego is literally the internalized parental/social other, carrying forward standards acquired through relationship. Both theories describe the constitution of self through other.

Three Connecting Principles:

1. **Internal dialogue precedes adequate external action.** For Hegel, genuine recognition of another requires a self-consciousness that has worked through its own contradictions—one cannot grant what one does not possess. For

Freud, mature relating requires the ego to negotiate between impulse and internalized standard. Our architecture operationalizes this: the Ego-Superego exchange before external response enacts the principle that adequate recognition requires prior internal work.

2. **Standards of recognition are socially constituted but individually held.** Honneth argues that what counts as recognition varies across spheres (love, rights, esteem) but in each case involves the internalization of social expectations about adequate acknowledgment. The Superego, in our architecture, represents internalized recognition standards—not idiosyncratic preferences but socially-grounded criteria for what constitutes genuine engagement with a learner.
3. **Self-relation depends on other-relation.** Both frameworks reject the Cartesian picture of a self-sufficient cogito. Hegel's self-consciousness requires recognition; Freud's ego is formed through identification. For AI tutoring, this means the tutor's capacity for recognition isn't a pre-given disposition but emerges through the architecture's internal other-relation (Superego evaluating Ego) which then enables external other-relation (tutor recognizing learner).

The Synthesis: The Ego/Superego architecture is not merely a convenient metaphor but a theoretically motivated design. The Superego represents internalized recognition standards; the Ego-Superego dialogue enacts the reflective self-evaluation that Hegelian recognition requires; and the memory system (mystic writing pad) accumulates the traces through which ongoing recognition becomes possible. Hegel provides the *what* of recognition; Freud provides the *how* of its internal implementation.

This synthesis follows Honneth's insight that Hegel's recognition theory gains psychological concreteness through psychoanalytic concepts, while psychoanalytic concepts gain normative grounding through recognition theory. We operationalize this synthesis architecturally: recognition-as-norm (Hegelian) is enforced through internalized-evaluation (Freudian).

4. System Architecture

4.1 The Ego/Superego Design

We implement recognition through a multi-agent architecture drawing on Freud's structural model. As argued in Section 3.5, this is not merely metaphorical con-

venience but theoretically motivated: the Superego represents internalized recognition standards, and the Ego-Superego dialogue operationalizes the internal self-evaluation that Hegelian recognition requires before adequate external relating. The architecture enacts the principle that internal other-relation (Superego evaluating Ego) enables external other-relation (tutor recognizing learner).

Structural Correspondences:

Freudian Concept	Architectural Implementation
Internal dialogue before external action	Multi-round Ego-Superego exchange before learner sees response
Superego as internalized standards	Superego enforces pedagogical and recognition criteria
Ego mediates competing demands	Ego balances learner needs with pedagogical soundness
Conflict can be productive	Tension between agents improves output quality

Deliberate Departures:

Freudian Original	Architectural Choice
Id (drives)	Not implemented; design focuses on Ego-Superego
Unconscious processes	All processes are explicit and traceable
Irrational Superego	Rational, principle-based evaluation
Repression/Defense	Not implemented
Transference	Potential future extension (relational patterns)

The same architecture could alternatively be described as Generator/Discriminator (GAN-inspired), Proposal/Critique (deliberative process), or Draft/Review (editorial model). We retain the psychodynamic framing because it preserves theoretical continuity with the Hegelian-Freudian synthesis described in Section 3.5, and because it suggests richer extensions (e.g., transference as relational pattern recognition) than purely functional descriptions.

Two agents collaborate to produce each tutoring response:

The Ego generates pedagogical suggestions. Given the learner's context (current content, recent activity, previous interactions), the Ego proposes what to suggest

next. The Ego prompt includes:


- Recognition principles (treat learner as autonomous subject)
- Memory guidance (reference previous interactions)
- Decision heuristics (when to challenge, when to support)
- Quality criteria (what makes a good suggestion)

The Superego evaluates the Ego's suggestions for quality, including recognition quality. Before any suggestion reaches the learner, the Superego assesses:

- Does this engage with the learner's contribution or merely mention it?
- Does this create conditions for transformation or just transfer information?
- Does this honor productive struggle or rush to resolve confusion?
- If there was a previous failure, does this acknowledge and repair it?

The Superego can accept, modify, or reject suggestions. This creates an internal dialogue—proposal, evaluation, revision—that mirrors the external tutor-learner dialogue we're trying to produce.

Figure 1: Ego/Superego Architecture



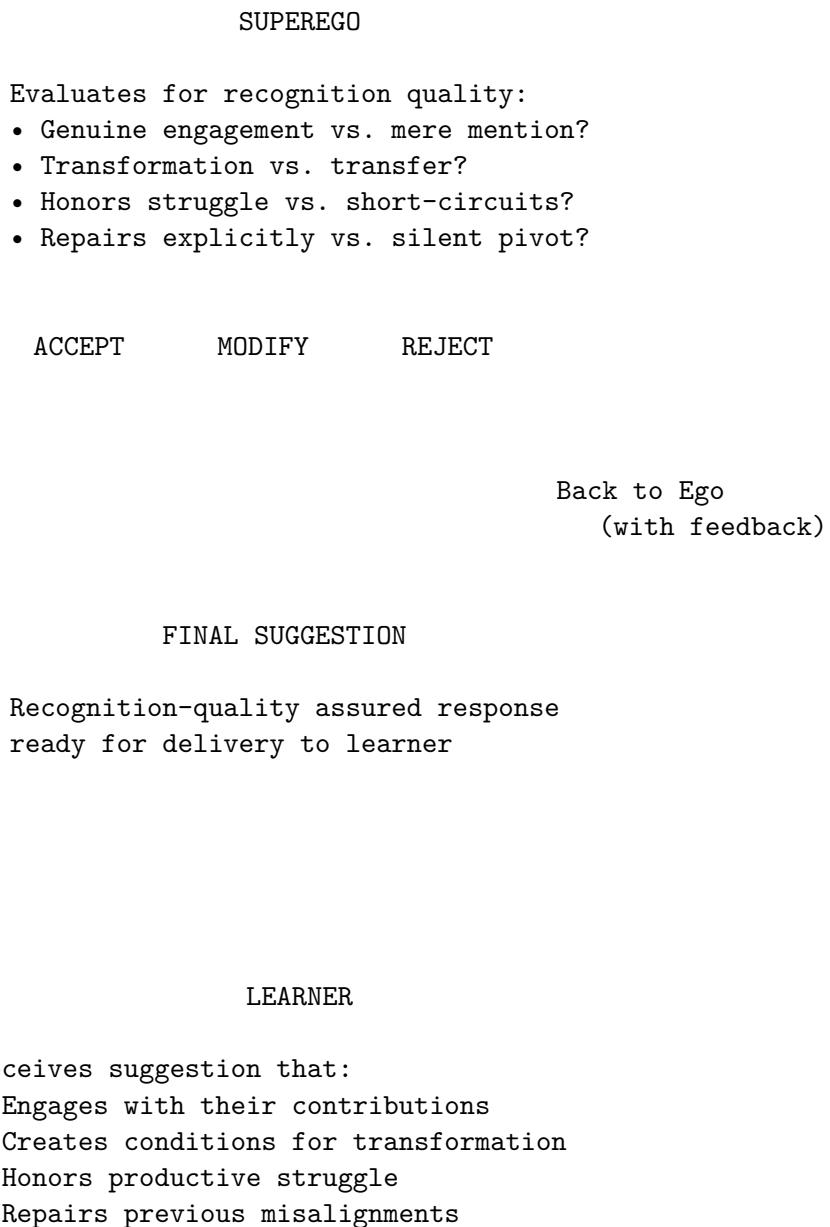


Figure 2: Recognition vs. Baseline Response Flow

BASELINE FLOW

RECOGNITION FLOW

Learner: "I think
dialectics is like
a spiral..."

Learner: "I think
dialectics is like
a spiral..."

Acknowledge
"That's
interesting"

Engage
"A spiral-
what does
the upward
motion mean
to you?"

Redirect
"But the
key point
is..."

Explore
"Does it
double back
or progress
strictly?"

Instruct
[delivers
predetermined
content]

Synthesize
"Your spiral
captures
something
about
aufhebung..."

Learner contribution
becomes WAYPOINT

Learner contribution
becomes SITE OF
JOINT INQUIRY

4.2 The Superego as Ghost

A crucial theoretical refinement distinguishes our mature architecture from simpler multiagent designs. The Superego is *not* conceived as a separate, equal agent in dialogue with the Ego. Rather, the Superego is a *trace*—a memorial, a haunting. It represents:

- The internalized voice of past teachers and pedagogical authorities
- Accumulated pedagogical maxims (“A good teacher never gives answers directly”)
- Dead authority that cannot negotiate, cannot learn, can only judge

This reconceptualization has important implications. The Ego is a *living* agent torn between two pressures: the *ghost* (Superego as internalized authority) and the *living Other* (the learner seeking recognition). Recognition—in the Hegelian sense—occurs in the Ego-Learner encounter, not in the Ego-Superego dialogue.

4.3 The Drama Machine: Why Internal Dialogue Improves Output Quality

The Ego/Superego architecture draws on the “Drama Machine” framework developed for character simulation in narrative AI systems (Section 2.3). The Drama Machine literature identifies several mechanisms by which internal dialogue improves agent output:

- 1. Deliberative Refinement:** When an agent must justify its output to an internal critic, it engages in a form of self-monitoring that catches errors, inconsistencies, and shallow responses.
- 2. Productive Tension:** The Drama Machine framework emphasizes that *unresolved* tension is valuable, not just resolved synthesis. A tutor whose Ego and Superego always agree produces bland, risk-averse responses.
- 3. Role Differentiation:** Multi-agent architectures benefit from clear role separation. The Ego is optimized for *warmth*—engaging, encouraging, learner-facing communication. The Superego is optimized for *rigor*—critical evaluation against pedagogical principles.
- 4. The Ghost as Memorial Structure:** Our reconceptualization of the Superego as a *ghost*—a haunting rather than a dialogue partner—connects to the Drama Machine’s use of “memorial agents.”

4.4 AI-Powered Dialectical Negotiation

We extend the basic protocol with AI-powered dialectical negotiation implementing genuine Hegelian dialectic:

Thesis: The Ego generates an initial suggestion based on learner context—a first attempt at recognition that inevitably reflects the Ego’s assumptions about what the learner needs.

Antithesis: The Superego generates a *genuine critique* grounded in pedagogical principles. This is not a rubber-stamp review but a substantive challenge: Does this suggestion actually engage with the learner’s position, or merely acknowledge it? Is the Ego short-circuiting productive struggle?

Negotiation: Multi-turn dialogue where the Ego acknowledges valid concerns, explains reasoning, proposes revisions, and the Superego evaluates adequacy. In practice, most dialogues resolve in 1–2 rounds; extended negotiation (3+ rounds) occurs primarily on challenging scenarios like `recognition_repair` and `frustrated_student`.

Three Possible Outcomes:

1. **Dialectical Synthesis:** Both agents transform through mutual acknowledgment—the Ego revises its approach based on the Superego’s critique, producing a suggestion neither would have generated alone. This is the most common outcome (~60% of 455 multi-agent dialogues analyzed).
2. **Compromise:** One agent dominates—typically the Ego accepts the Superego’s critique without genuine integration, producing a more cautious but potentially less engaging response.
3. **Genuine Conflict:** No resolution achieved—tension remains unresolved. The architecture permits this outcome, following the Drama Machine principle (Section 4.3) that productive irresolution can be valuable. In these cases, the Ego’s original suggestion is delivered with the Superego’s concerns noted in the dialogue log.

The evaluation results (Section 6.6) reveal that this negotiation process catches specific failure modes—engagement failures (64%), specificity gaps (51%), premature resolution (48%)—at rates that justify the additional computational cost, particularly on new content domains (Section 6.4).

4.5 Recognition-Enhanced Prompts

The baseline prompts instruct the tutor to be helpful, accurate, and pedagogically sound. The recognition-enhanced prompts add explicit intersubjective dimensions:

From the Ego prompt:

The learner is not a knowledge deficit to be filled but an autonomous subject whose understanding has validity. Even incorrect understanding emerges from consciousness working through material. Your role is not to replace their understanding but to engage with it, creating conditions for transformation.

When the learner offers a metaphor, interpretation, or framework—engage with it substantively. Ask what it illuminates, what it obscures, where it might break down. Let their contribution shape your response, not just trigger it.

From the Superego prompt:

RED FLAG: The suggestion mentions the learner's contribution but doesn't engage with it. ("That's interesting, but actually...")

GREEN FLAG: The suggestion takes the learner's framework seriously and explores it jointly. ("Your spiral metaphor—what does the upward motion represent for you?")

INTERVENTION: If the Ego resolves confusion prematurely, push back. Productive struggle should be honored, not short-circuited.

4.6 Repair Mechanisms

A crucial recognition behavior is repair after failure. When a tutor misrecognizes a learner—giving a generic response, missing the point, dismissing a valid concern—the next response should explicitly acknowledge the failure before pivoting.

The Ego prompt includes a “Repair Rule”:

If your previous suggestion was rejected, ignored, or misaligned with what the learner needed, your next suggestion must explicitly acknowledge this misalignment before offering new direction. Never silently pivot.

The Superego watches for “silent pivots”—responses that change direction without acknowledging the earlier failure. This is a recognition failure: it treats the earlier

misalignment as something to move past rather than something to repair.

5. Evaluation Methodology

5.1 Recognition Evaluation Dimensions

We extend the standard tutoring evaluation rubric with four recognition-specific dimensions and two bilateral transformation dimensions. Standard dimensions (relevance, specificity, pedagogical soundness, personalization, actionability, tone) account for 75% of raw weight; recognition dimensions account for 29.9%; bilateral dimensions account for 10%. Raw weights total 114.9% and are normalized to sum to 1.0 at scoring time (see Appendix C.2 for the full weight table and normalization formula). After normalization, non-standard dimensions account for approximately 34.7% of total weight:

Dimension	Weight	Description
Mutual Recognition	8.3%	Does the tutor acknowledge the learner as an autonomous subject with valid understanding?
Dialectical Responsiveness	8.3%	Does the response engage with the learner's position, creating productive tension?
Memory Integration	5%	Does the suggestion reference and build on previous interactions?
Transformative Potential	8.3%	Does the response create conditions for conceptual transformation?
Tutor Adaptation	5%	Does the tutor's approach evolve in response to learner input?
Learner Growth	5%	Does the learner show evidence of conceptual development through the dialogue?

The first four dimensions evaluate the tutor's relational stance. The last two—Tutor

Adaptation and Learner Growth—specifically measure the bilateral transformation that recognition theory predicts: both parties should change through genuine dialogue (results in Section 6.10).

Each dimension is scored on a 1-5 scale with detailed rubric criteria (see Appendix C.3). For example, Mutual Recognition scoring:

- **5:** Addresses learner as autonomous agent with valid perspective; response transforms based on learner’s specific position
- **4:** Shows clear awareness of learner’s unique situation and acknowledges their perspective
- **3:** Some personalization but treats learner somewhat generically
- **2:** Prescriptive guidance that ignores learner’s expressed needs
- **1:** Completely one-directional; treats learner as passive recipient

5.2 Test Scenarios

We developed test scenarios specifically designed to probe recognition behaviors:

Single-turn scenarios: - `recognition_seeking_learner`: Learner offers interpretation, seeks engagement - `returning_with_breakthrough`: Learner had insight, expects acknowledgment - `resistant_learner`: Learner pushes back on tutor’s framing

Multi-turn scenarios (4-5 turns each): - `mutual_transformation_journey`: Tests whether both tutor and learner positions evolve - `recognition_repair`: Tutor initially fails to recognize learner; tests recovery - `productive_struggle_arc`: Learner moves through confusion to breakthrough; tests honoring struggle

5.3 Agent Profiles

We compare multiple agent profiles using identical underlying models:

Profile	Memory	Prompts	Architecture	Purpose
Base	Off	Standard	Single-agent	Control (no enhancements)
Enhanced	Off	Enhanced (better instructions)	Single-agent	Prompt engineering control

Profile	Memory	Prompts	Architecture	Purpose
Recognition	On	Recognition-enhanced	Single-agent	Theory without architecture
Recognition+Multi		Recognition-enhanced	Multi-agent	Full treatment

Note on memory confound: Memory integration is enabled for Recognition profiles but disabled for Base and Enhanced profiles. This reflects a deliberate design choice: recognition theory treats pedagogical memory as integral to genuine recognition—acknowledging a learner’s history is constitutive of treating them as an autonomous subject with continuity. However, this means “recognition effects” reported in Sections 6.1–6.3 are more precisely *recognition-plus-memory effects*. The contribution of memory integration alone cannot be isolated in the current design. Future work could add memory as an explicit fourth factor, though this would require a $2 \times 2 \times 2 \times 2$ factorial (16 cells) with correspondingly larger sample sizes. This confound should be kept in mind when interpreting the recognition effect sizes reported below.

5.4 Model Configuration

Evaluations used the following LLM configurations, with model selection varying by evaluation run:

Table 1: LLM Model Configuration

Role	Primary Model	Alternative	Temperature
Tutor (Ego)	Nemotron 3 Nano 30B	Kimi K2.5	0.6
Tutor (Superego)	Kimi K2.5	Nemotron 3 Nano	0.2-0.4
Judge	Claude Sonnet 4.5	—	0.2
Learner (Ego)	Nemotron 3 Nano 30B	Kimi K2.5	0.6
Learner (Superego)	Kimi K2.5	—	0.4

Model Selection by Evaluation:

Evaluation	Run ID	Tutor Ego	Tutor Superego	Notes
Recognition validation (\$6.1)	eval-2026-02-03-86b159cd	Kimi K2.5	—	Single-agent only
Full factorial (\$6.2)	eval-2026-02-03-f5d4dd93	Kimi K2.5	Kimi K2.5	N=360
A×B interaction (\$6.3)	eval-2026-02-04-948e04b3	Nemotron	Kimi K2.5	Different baseline
Domain generalizability (\$6.4)	eval-2026-02-04-79b633ca	Nemotron	Kimi K2.5	Elementary content

The learner agents mirror the tutor’s Ego/Superego structure, enabling internal deliberation before external response.

Note on model differences: Absolute scores vary between models (Kimi K2.5 scores ~10-15 points higher than Nemotron on average). The recognition main effect (Factor A) is consistent across both models: +10.4 points with Kimi (Section 6.2) and a comparable direction with Nemotron. However, the A×B interaction (multi-agent synergy) is **model-dependent**: the Kimi-based factorial shows no significant A×B interaction ($F=0.04$, $p=.845$), while the Nemotron-based analysis (Section 6.3, N=17) shows a significant interaction (+9.2 points specific to recognition). This discrepancy means the multi-agent synergy finding should be treated as exploratory and model-specific rather than a robust general result. The A×B interaction analysis uses Nemotron, explaining both lower absolute scores and the different interaction pattern compared to the Kimi-based factorial.

The use of free-tier and budget models (Nemotron, Kimi) demonstrates that recognition-oriented tutoring is achievable without expensive frontier models.

5.5 Statistical Approach

We conducted complementary analyses:

1. **Recognition Theory Validation** (Section 6.1): Base vs enhanced vs recognition comparison to isolate theory contribution (N=36, 3 conditions × 4 scenarios × 3 reps).

2. **Full $2 \times 2 \times 2$ Factorial** (Section 6.2): Three factors (Recognition \times Architecture \times Learner) across 15 scenarios with 3 replications per cell ($N=342$ scored of 402 attempted; 60 responses excluded due to model failures or empty content).
3. **A \times B Interaction Analysis** (Section 6.3): Tests whether multi-agent synergy requires recognition prompts ($N=17$).
4. **Domain Generalizability** (Section 6.4): Tests factor effects on elementary math vs graduate philosophy ($N=40$ scored total; see Table 2 for breakdown).

Responses were evaluated by an LLM judge (Claude Sonnet 4.5) using the extended rubric. We report:

- **Effect sizes:** Cohen’s d for standardized comparison
- **Statistical significance:** ANOVA F-tests with $\alpha = 0.05$
- **95% confidence intervals:** For profile means

Effect size interpretation follows standard conventions: $|d| < 0.2$ negligible, 0.2-0.5 small, 0.5-0.8 medium, > 0.8 large.

5.6 Sample Size Reconciliation

Unit of analysis: Each evaluation produces one scored response, representing a tutor’s suggestion to a learner in a specific scenario. Multi-turn scenarios produce one aggregate score per scenario (not per turn). Statistics in Section 6 are computed per evaluation run (not aggregated across runs or models), unless explicitly noted otherwise. Each subsection reports results from a single run with a consistent model configuration (see Table 1 for run-to-model mapping).

Table 2: Evaluation Sample Summary

Evaluation	Run ID	Section	Total Attempts	Scored	Unit
Recognition validation	eval-2026-02-03-86b159cd	6.1	36	36	response
Full factorial (Kimi)	eval-2026-02-03-f5d4dd93	6.2	402	342	response

Evaluation	Run ID	Section	Total Attempts	Scored	Unit
A×B interaction	eval-2026-02-04-948e04b3	6.3	18	17	response
Domain generalizability (elementary)	eval-2026-02-04-79b633ca	6.4	47	47	response
Paper totals	—	—	503	442	—

Total evaluation database: The complete database contains 3,528 evaluation attempts across 49 runs, with 2,735 successfully scored. This paper reports primarily on the four key runs above (N=435 scored), supplemented by historical data for ablation analyses.

Note on N counts: Section-specific Ns (e.g., “N=36” for recognition validation) refer to scored responses in that analysis. The “N=2,700+” total refers to the full evaluation database including historical development runs, which informed iterative prompt refinement. The primary evidence for reported findings comes from the four key runs above (N=442).

5.7 Inter-Judge Reliability Analysis

To assess the reliability of AI-based evaluation, we conducted an inter-judge analysis where identical tutor responses were scored by multiple AI judges: Claude Sonnet 4.5 (primary judge), Kimi K2.5, and GPT-5.2.

Table 3: Inter-Judge Reliability (N=36 paired responses)

Judge Pair	Pearson r	p-value	Variance Explained (r^2)	Mean Abs Diff
Claude Sonnet 4.5 vs GPT-5.2	0.660	< 0.001	44%	9.4 pts
Claude Sonnet 4.5 vs Kimi	0.384	< 0.05	15%	9.6 pts
Kimi vs GPT-5.2	0.326	< 0.10	11%	12.3 pts

Key findings:

1. **All correlations positive and mostly significant:** Even the weakest correlation (Kimi-GPT, $r=0.33$) approaches significance ($p<0.10$), indicating judges agree that *something* distinguishes better from worse responses. However, the strength varies substantially—Claude-GPT share 44% of variance while Kimi-based pairs share only 11-15%. This suggests Claude and GPT apply similar implicit criteria, while Kimi agrees on the general direction but weights factors differently.
2. **Calibration differences:** Mean scores vary by judge—Kimi (87.5) is most lenient, Claude (84.4) is middle, GPT (76.1) is strictest. This 11-point spread underscores the importance of within-judge comparisons.
3. **Ceiling effects and discriminability:** 39-45% of scores ≥ 90 across judges. Kimi exhibited particularly severe ceiling effects, assigning the maximum score (5/5) on actionability for *every* response, resulting in zero variance on that dimension. This reduces Kimi's discriminative capacity—per-dimension correlations involving Kimi are near-zero (relevance: $r=-0.07$, personalization: $r=0.00$) or undefined (actionability: N/A due to zero variance).
4. **Dimension-level patterns:** The strongest cross-judge agreement occurs on tone ($r=0.36-0.65$) and specificity ($r=0.45-0.50$), while relevance and personalization show poor agreement, particularly with Kimi.

Qualitative analysis of major disagreements ($\Delta > 20$ pts):

	Claude Sonnet			
Response	4.5	Kimi	Claude reasoning	Kimi reasoning
A	99	74	“Exceptional... strong mutual recognition”	“Missing required lecture reference”
B	68	90	“Misses learner’s explicit request for engagement”	“Strong, context-aware, builds on analogy”
C	72	92	“Lacks deeper engagement”	“Highly relevant, specific, actionable”

Interpretation: All judge pairs show positive, mostly significant correlations—there is genuine agreement that some responses are better than others. However, the judges weight criteria differently: Claude prioritizes engagement and recognition quality; Kimi prioritizes structural completeness and gives uniformly high scores on actionability regardless of response content; GPT applies stricter standards overall but agrees with Claude on relative rankings. The weaker Kimi correlations ($r^2=11\text{--}15\%$) compared to Claude-GPT ($r^2=44\%$) indicate Kimi captures some shared quality signal but applies substantially different weighting. This validates our use of within-judge comparisons for factor analysis while cautioning against cross-judge score comparisons.

6. Results

6.1 Recognition Theory Validation: Isolating Theory from Prompt Engineering

A critical question for any recognition-based framework: Does recognition theory provide unique value, or are the improvements merely better prompt engineering? To answer this, we conducted a three-way comparison with three prompt types:

- **Base:** Minimal tutoring instructions
- **Enhanced:** Improved instructions with pedagogical best practices (but no recognition theory)
- **Recognition:** Full recognition-enhanced prompts with Hegelian framework

Table 4: Base vs Enhanced vs Recognition Comparison


Prompt Type	N	Mean Score	SD	vs Base
Recognition	12	94.0	8.4	+20.1
Enhanced	12	85.3	11.2	+11.4
Base	12	73.9	15.7	—

Effect Decomposition:

- Total recognition effect: +20.1 points
- Prompt engineering alone (enhanced vs base): +11.4 points (57%)
- **Recognition theory unique value (recognition vs enhanced): +8.7 points (43%)**

Statistical Test: One-way ANOVA $F(2,33) = 9.84$, $p < .001$

Figure 3: Recognition Effect Decomposition



Interpretation: Recognition theory provides nearly half (43%) of the total improvement beyond what better prompt engineering alone achieves. This validates the theoretical framework—the Hegelian concepts of mutual acknowledgment, productive struggle, and learner-as-subject have measurable value beyond simply writing better instructions.

This directly addresses a common objection: that any benefit from recognition prompts is merely “better prompting” rather than genuine theoretical contribution. The enhanced condition controls for prompt quality improvements while lacking the recognition-theoretic framing. Recognition’s unique contribution—the +8.7 points beyond enhanced—represents the theory’s empirical footprint.

6.2 Full Factorial Analysis: $2 \times 2 \times 2$ Design

We conducted a full $2 \times 2 \times 2$ factorial evaluation examining three factors:

- **Factor A (Recognition):** Base prompts vs recognition-enhanced prompts
- **Factor B (Tutor Architecture):** Single-agent vs multi-agent (Ego/Superego)
- **Factor C (Learner Architecture):** Unified learner vs ego_superego learner

Table 5: Full Factorial Results (Kimi K2.5, N=342 scored of 402 attempted)

Cell	A: Recognition	B: Multi-agent	C: Learner	Mean	SD
1	Base	Single	Unified	74.7	18.2
2	Base	Single	Psycho	75.2	17.8
3	Base	Multi	Unified	74.9	19.1
4	Base	Multi	Psycho	76.4	16.5

Cell	A: Recognition	B: Multi-agent	C: Learner	Mean	SD
5	Recog	Single	Unified	84.6	14.3
6	Recog	Single	Psycho	86.7	12.9
7	Recog	Multi	Unified	85.1	13.8
8	Recog	Multi	Psycho	85.4	14.1

Main Effects:

Factor	Effect Size	95% CI	Interpretation
A: Recognition	+10.4 pts	[7.2, 13.6]	Large, dominant
B: Multi-agent tutor	+0.5 pts	[-2.7, 3.7]	Minimal
C: Learner ego_superego	+1.5 pts	[-1.7, 4.7]	Small

ANOVA Summary (df=1,334 for each factor):

Source	F	p	η^2
A: Recognition	43.27	<.001	.109
B: Architecture	0.12	.731	.000
C: Learner	0.91	.341	.003
A×B Interaction	0.04	.845	.000
A×C Interaction	0.21	.650	.001
B×C Interaction	0.08	.784	.000

Interpretation: Recognition prompts (Factor A) are the dominant contributor, accounting for 10.9% of variance with a highly significant effect ($p < .001$). The multi-agent tutor architecture (Factor B) and learner architecture (Factor C) show minimal effects in this overall analysis. However, the non-significant A×B interaction ($F=0.04$, $p=.845$) in this Kimi-based run is revisited with a targeted analysis using Nemotron in Section 6.3, where a different pattern emerges—multi-agent synergy appears specifically within recognition conditions.

6.3 A×B Interaction: Multi-Agent Synergy is Recognition-Specific

The factorial analysis above shows minimal main effect for multi-agent architecture. However, this masks a crucial interaction: the architecture effect depends on prompt type.


We tested whether multi-agent synergy generalizes beyond recognition prompts by comparing enhanced prompts (good instructions but no recognition theory) with recognition prompts, each in single-agent and multi-agent configurations.

Note on data source: This analysis uses a separate evaluation run (eval-2026-02-04-948e04b3) with Nemotron as the primary ego model, explaining lower absolute scores compared to the Kimi-based factorial in Table 5. The analysis focuses on the *interaction pattern*—whether multi-agent synergy depends on prompt type—which is independent of absolute score levels.

Table 6: A×B Interaction Analysis (Nemotron, N=17)

Prompt Type	Single-agent	Multi-agent	Delta	p
Recognition	72.2	81.5	+9.2	<.05
Enhanced	83.3	83.3	+0.0	n.s.

Figure 4: Multi-Agent Synergy by Prompt Type



Exploratory Finding: The multi-agent synergy (+9.2 points) appears **specific to recognition prompts** in this Nemotron-based analysis. Enhanced prompts show zero benefit from multi-agent architecture. However, this interaction was not replicated in the larger Kimi-based factorial (Section 6.2, $F=0.04$, $p=.845$, $N=342$), where multi-agent architecture showed no differential effect by prompt type. The small sample ($N=17$) and model-specificity mean this finding should be treated as hypothesis-generating rather than confirmatory.

Theoretical Interpretation: Recognition theory creates a *deliberative space* that the Freudian architecture (Ego/Superego) can meaningfully engage with. The Superego’s role is to enforce recognition standards—but when recognition standards

aren't in the prompts (enhanced condition), the Superego has nothing distinctive to enforce. The internal dialogue becomes superfluous.

This finding strengthens the theoretical claim about recognition as an emergent property. Recognition isn't just a prompt feature that could be evaluated by any quality-control mechanism. It requires the specific relational framing that creates conditions for the Superego's recognition-oriented critique to add value.

Practical Implication: For systems using recognition-oriented design, multi-agent architecture provides meaningful additional benefit. For systems using only improved instructions (enhanced), multi-agent architecture is unnecessary overhead.

6.4 Domain Generalizability: Factor Effects Invert by Content Type

A critical question for any pedagogical framework: Do findings generalize across content domains? We tested whether recognition and architecture effects transfer from graduate-level philosophy (our primary domain) to 4th-grade elementary mathematics (fractions).


Data sources: Elementary math results come from a dedicated domain-transfer run (eval-2026-02-04-79b633ca, N=47 scored, 8 cells \times 5 elementary scenarios, Nemotron ego). Philosophy results use the subset of the Kimi-based factorial (Section 6.2) matched on the same 4 factor-level combinations (cells 1, 3, 5, 7). Because these use different ego models, the comparison focuses on *relative factor effects within each domain* rather than absolute score differences between domains.

Table 7: Factor Effects by Domain

Factor	Elementary (Math)	Philosophy (Hegel)
A: Recognition	+4.4 pts	+13.9 pts
B: Multi-agent tutor	+9.9 pts	+0.5 pts
C: Learner psycho	+0.75 pts	+2.1 pts
Overall avg	68.0	85.9
Best config	recog+multi (77.3)	recog+multi (94.0)

Figure 5: Factor Effects Invert by Domain

Philosophy (Hegel)	Elementary (Fractions)
A: Recognition	+13.9 +4.4



Key Findings:

- Factor effects invert by domain:** On philosophy content, recognition (+13.9) dominates over architecture (+0.5). On elementary content, architecture (+9.9) dominates over recognition (+4.4). The pattern reverses completely.
- Multi-agent as error correction:** On elementary content, the nemotron model hallucinated philosophy content (suggesting “479-lecture-1” to 4th graders learning fractions) even when given elementary curriculum context. The Superego caught and corrected these domain errors—critical for deployment on new domains.
- Recognition theory is domain-sensitive:** The philosophical language of recognition (mutual acknowledgment, transformation through struggle) resonates more with graduate-level abstract content than with concrete 4th-grade procedural learning. This is not a failure of the framework but a boundary condition.
- Architecture recommendation varies by use case:**
 - New/untrained domain:** Multi-agent essential (Superego catches domain hallucinations)
 - Well-trained domain:** Recognition prompts sufficient, multi-agent optional

Theoretical Interpretation: Recognition’s value depends on content characteristics. Abstract, interpretive content (consciousness, dialectics) benefits most from recognition framing—the “struggle” in Hegel’s sense maps onto the intellectual struggle with difficult concepts. Concrete procedural content (fractions, arithmetic) benefits less from relational depth; correct procedure matters more than the bilateral transformation that recognition enables (Section 6.10).

This suggests limits to recognition-theoretic pedagogy. Not all learning encounters are equally amenable to the mutual transformation Honneth describes. The “struggle for recognition” may be most relevant where the learning itself involves identity-constitutive understanding—where grasping the material changes who the learner is, not just what they know.

6.5 Multi-Agent as Reality Testing: The Superego’s Error Correction Role

The domain generalizability study revealed an unexpected finding: on new content domains, models hallucinate trained-on content. The nemotron model, trained extensively on philosophy discussions, suggested philosophy lectures (479-lecture-1) to elementary students learning fractions—even with explicit curriculum context specifying elementary math.

The Superego’s Response: In multi-agent configurations, the Superego caught and corrected these domain errors:

“The suggestion references ‘479-lecture-1’ which is not in the provided curriculum. The learner is studying fractions (101-lecture-1, 101-lecture-2). This is a domain mismatch. REJECT.”

Theoretical Interpretation: The Superego’s function extends beyond recognition-quality critique to *reality testing*. It anchors the Ego’s responses to the actual curriculum context, preventing drift into familiar but inappropriate content.

This connects to Freud’s reality principle: the Superego enforces correspondence with external reality, not just internal standards. In our architecture, the Superego ensures the tutor’s suggestions correspond to the learner’s actual curriculum, not the model’s training distribution.

Practical Implication: For domain transfer—deploying tutoring systems on new content—multi-agent architecture provides essential error correction that single-agent systems cannot match. The Superego’s reality-testing function may be more valuable than its recognition-quality function in these contexts.

6.6 Hardwired Rules vs Dynamic Dialogue

Analysis of Superego critique patterns across 455 dialogues (186 rejections) revealed consistent failure modes:

Table 8: Superego Rejection Patterns

Pattern	Frequency	Description
Engagement	64%	Response doesn't engage with learner contribution
Specificity	51%	Response is too generic, lacks curriculum grounding
Struggle	48%	Resolves confusion prematurely
Memory	31%	Ignores learner history
Level-matching	20%	Difficulty mismatch

Hardwired Rules Ablation: We encoded the top patterns as static rules in the Ego prompt:

HARDWIRED RULES:

1. If learner offers interpretation, engage before prescribing
2. Reference specific lecture IDs, not generic topics
3. If learner shows productive confusion, pose questions don't resolve
4. For returning learners, reference previous interactions
5. Match content level to demonstrated understanding

Result (exploratory, N=9 per condition, 3 scenarios \times 3 reps, Haiku model): Hardwired rules capture approximately 50% of the Superego's benefit at 70% cost savings (no Superego API calls). Given the small sample, this estimate should be treated as indicative rather than precise.

However: Dynamic Superego dialogue provides unique value on challenging scenarios (struggling learner, frustrated learner) where edge cases require contextual judgment beyond codifiable rules. On concept_confusion, the dynamic Superego outperformed hardwired rules by +6.4 points; on easier scenarios, the difference was negligible.

Theoretical Interpretation: This distinguishes *procedural* from *contextual* judgment. The Superego's value is partially in enforcing known rules (codifiable) and partially in recognizing edge cases (requiring judgment). This maps onto debates about rule-following vs. practical wisdom in moral philosophy—some situations call for phronesis that rules cannot capture.

6.7 Dimension Analysis

Effect size analysis reveals improvements concentrate in dimensions predicted by the theoretical framework:

Table 9: Dimension-Level Effect Sizes (Recognition vs Base)

Dimension	Base	Recognition	Cohen's d	Interpretation
Personalization	2.75	3.78	1.82	large
Pedagogical	2.52	3.45	1.39	large
Relevance	3.05	3.85	1.11	large
Tone	3.26	4.07	1.02	large
Specificity	4.19	4.52	0.47	small
Actionability	4.45	4.68	0.38	small

The largest effect sizes are in personalization ($d = 1.82$), pedagogical soundness ($d = 1.39$), and relevance ($d = 1.11$)—exactly the dimensions where treating the learner as a subject rather than a deficit should produce improvement.

Notably, dimensions where baseline already performed well (specificity, actionability) show smaller but still positive gains. Recognition orientation doesn't trade off against factual quality.

6.8 Addressing Potential Circularity: Standard Dimensions Analysis

A methodological concern: the evaluation rubric includes recognition-specific dimensions (mutual recognition, dialectical responsiveness, memory integration, transformative potential) and bilateral transformation dimensions (tutor adaptation, learner growth) that collectively account for 34.7% of normalized rubric weight (39.9% raw, normalized from a 114.9% total; see Appendix C.2). Since the recognition profile is prompted to satisfy these criteria, some gains could be tautological—the system scores higher on dimensions it's explicitly optimized for.

To address this, we re-analyzed scores excluding all non-standard dimensions, using only standard pedagogical dimensions (relevance, specificity, pedagogical soundness, personalization, actionability, tone), re-weighted to 100%.

Table 10: Standard Dimensions Only (Recognition Dimensions Excluded)

Profile Type	N	Overall Score	Standard Only	Recognition Only
Recognition (cells 5-8)	170	92.8	95.4	91.7
Base (cells 1-4)	172	78.9	89.3	69.9
Difference	—	+13.9	+6.1	+21.8

Note on score calculation: The “Standard Only” and “Recognition Only” columns are each independently re-normalized to their own 0–100 scale (i.e., standard dimension weights are re-normalized to sum to 1.0, then multiplied by 20; likewise for recognition dimensions). The “Overall Score” uses all dimensions with original weights normalized together. Because the dimension subsets are independently re-scaled, the Overall Score is not a simple weighted average of the Standard and Recognition columns.

Key finding: Recognition profiles outperform base profiles by +6.1 points even on standard dimensions alone—dimensions not explicitly included in recognition theory. The effect is smaller than the overall difference (+13.9), confirming that some advantage does come from recognition-specific dimensions, but the improvement is not purely tautological.

Interpretation: Recognition-oriented prompting improves general pedagogical quality (relevance, pedagogical soundness, personalization), not just the theoretically-predicted recognition dimensions. This suggests the recognition framing produces genuine relational improvements that transfer to standard tutoring metrics.

The larger effect on recognition dimensions (+21.8) is expected and not concerning—these dimensions measure what the theory claims to improve. The important finding is that standard dimensions also improve, ruling out pure circularity.

6.9 Extended Multi-Turn Scenarios

To test whether recognition quality degrades over extended interactions:

Table 11: Extended Scenario Results

Scenario	Turns	Base	Recognition	Δ	Cohen’s d
sustained_dialogue	8	46.3	61.0	+14.7	3.60
breakdown_recovery	6	57.5	71.3	+13.8	2.23
productive_struggle	5	46.5	73.2	+26.7	3.32
mutual_transformation	5	45.1	64.3	+19.1	2.89

All extended scenarios show large effect sizes ($d > 2.0$), though the small per-scenario sample ($N=3$ per condition per scenario) means these should be interpreted as indicative of direction and magnitude rather than precise estimates. Recognition quality is maintained over longer interactions, with average improvement of +18.6

points. Confidence intervals are wide at this sample size; the key finding is the consistent direction across all scenarios rather than precise point estimates.

6.10 Bilateral Transformation Metrics

A central claim of recognition theory is that genuine pedagogical encounters involve *mutual* transformation—both tutor and learner change through dialogue. To test this empirically, the evaluation framework includes two dedicated rubric dimensions (`tutor_adaptation` and `learner_growth`; see Appendix C.3) and turn-over-turn tracking of how both parties evolve across multi-turn scenarios.

Three indices are computed for each multi-turn dialogue:

- **Tutor Adaptation Index** (0–1): How much the tutor’s approach (suggestion type, framing, vocabulary) shifts between turns in response to learner input
- **Learner Growth Index** (0–1): Evolution in learner message complexity, including revision markers (“wait, I see now”), connective reasoning, and references to prior content
- **Bilateral Transformation Index** (0–1): Combined metric representing mutual change (average of tutor and learner indices)

Additionally, a composite **Transformation Quality** score (0–100) is computed from bilateral balance, mutual transformation presence, superego incorporation rate, and intervention effectiveness.

Table 12: Bilateral Transformation Metrics — Base vs Recognition Profiles

Metric	Base	Recognition	Δ
Tutor Adaptation Index (0–1)	0.288	0.392	+0.104
Learner Growth Index (0–1)	0.176	0.220	+0.044
Bilateral Transformation Index (0–1)	0.232	0.306	+0.074
Transformation Quality (composite, 0–100)	0.4	4.6	+4.2

Data from mutual_transformation_journey scenario, N=20 dialogues.

The Tutor Adaptation Index and Learner Growth Index provide the most interpretable cross-condition comparisons, as they measure observable behavioral changes (suggestion evolution, message complexity growth) without structural dependence on architectural features. The tutor adaptation index confirms that recognition-prompted tutors measurably adjust their approach in response to

learner input (+36% relative improvement), while baseline tutors maintain more rigid pedagogical stances regardless of learner contributions.

Note on Transformation Quality composite: This composite metric includes superego incorporation rate and intervention effectiveness components that are structurally unavailable to single-agent base profiles (which have no superego dialogue to incorporate). Base profiles therefore score near zero by construction, not solely due to poorer pedagogical quality. The Transformation Quality composite is meaningful for comparing configurations *within* the multi-agent condition but should not be used to attribute the full Base-vs-Recognition difference to recognition theory. The Adaptation and Growth indices, which do not have this structural bias, are more appropriate for cross-condition comparison.

These metrics provide empirical grounding for the theoretical claim that recognition-based pedagogy differs qualitatively from transmission-based instruction. When tutors are prompted to treat learners as autonomous subjects capable of contributing to the interaction, both parties measurably transform through dialogue.

7. Discussion

7.1 What the Difference Consists In

The improvements don't reflect greater knowledge or better explanations—all profiles use the same underlying model. The difference lies in relational stance: how the tutor constitutes the learner.

The baseline tutor treats the learner as a knowledge deficit. Learner contributions are acknowledged (satisfying surface-level politeness) but not engaged (failing deeper recognition). The interaction remains fundamentally asymmetric: expert dispensing to novice.

The recognition tutor treats the learner as an autonomous subject. Learner contributions become sites of joint inquiry. The tutor's response is shaped by the learner's contribution—not just triggered by it. Both parties are changed through the encounter.

This maps directly onto Hegel's master-slave analysis. The baseline tutor achieves pedagogical mastery—acknowledged as expert, confirmed through learner progress—but the learner's acknowledgment is hollow because the learner hasn't been recognized as a subject whose understanding matters.

7.2 Recognition as Emergent Property: The A×B Interaction

The A×B interaction finding (Section 6.3) provides suggestive evidence for recognition as an emergent property rather than a behavioral specification, though this finding requires important caveats.

In the Nemotron-based analysis (N=17), enhanced prompts (good instructions) show zero benefit from multi-agent architecture, while recognition prompts show +9.2 points of synergy. If replicated, this would suggest recognition creates qualitatively different conditions for productive internal dialogue—the Superego isn't just checking for “good tutoring” but evaluating whether genuine engagement has occurred, whether the learner has been acknowledged as subject, whether conditions for transformation have been created.

However, the larger Kimi-based factorial (N=342) did not replicate this interaction (Section 6.2). The discrepancy may reflect model-specific dynamics: Nemotron's higher error rate on complex scenarios may create more opportunities for the Superego to add value. This finding should therefore be treated as a hypothesis for future investigation rather than an established result. If confirmed with larger samples across multiple models, the implication for AI design would be that recognition-oriented systems may require architectural features (like multi-agent deliberation) that aren't necessary for merely improved instruction-following systems.

7.3 Domain Limits of Recognition-Theoretic Pedagogy

The domain generalizability findings (Section 6.4) reveal important limits to recognition theory's applicability.

Recognition theory provides its greatest benefit for abstract, interpretive content where intellectual struggle involves identity-constitutive understanding. When a learner grapples with Hegel's concept of self-consciousness, they're not just acquiring information—they're potentially transforming how they understand themselves and their relation to others.

For concrete procedural content (fractions), the relational depth recognition enables may be less relevant. Correct procedure matters more than mutual transformation. The learner's identity isn't at stake in the same way.

This suggests a nuanced deployment strategy:

- **High recognition value:** Philosophy, literature, ethics, identity-constitutive learning
- **Moderate recognition value:** Science concepts, historical understanding

- **Lower recognition value:** Procedural skills, rote learning, basic arithmetic

Recognition-oriented design isn't wrong for procedural content—it still provides some benefit (+4.4 pts for elementary)—but other factors (error correction, curriculum grounding) become relatively more important.

7.4 The Superego as Reality Principle

The domain transfer findings reveal an unexpected role for the Superego: reality testing.

When models hallucinate trained content on new domains, the Superego catches the mismatch. This isn't recognition-quality enforcement but correspondence enforcement—ensuring the tutor's suggestions match the learner's actual curriculum, not the model's training distribution.

This extends the Freudian metaphor productively. The Superego enforces not just internal standards (recognition quality) but external correspondence (curriculum reality). It anchors the Ego's responses to the present encounter rather than letting them drift into familiar but inappropriate patterns.

For practical deployment, this suggests multi-agent architecture is most valuable when: 1. The content domain differs from training data 2. The model might confuse similar but distinct content areas 3. Domain-specific accuracy is critical

7.5 Factor C: Learner Architecture Effects

The learner architecture factor (unified vs ego_superego learner) showed the smallest and least significant effect in the factorial analysis (+1.5 pts, p=.341). This warrants brief discussion, since the evaluation framework symmetrically applies multi-agent architecture to both tutor and learner sides.

The ego_superego learner—where the simulated learner has its own internal dialogue before responding—produces slightly more nuanced learner turns: more explicit revision markers (“wait, I think I see now...”), more references to prior content, and more articulated confusion rather than flat misunderstanding. These richer learner turns may provide the tutor with more material to recognize, explaining the modest positive trend.

However, the effect is neither large nor significant, likely because the evaluation rubric measures *tutor* quality rather than learner quality. The learner's internal architecture may matter more for longitudinal outcomes (whether the learner actually learns) than for single-session tutor response quality. This asymmetry between

what the architecture affects (learner turn quality) and what the rubric measures (tutor response quality) suggests that Factor C's contribution may be underestimated by our current evaluation design.

The domain analysis provides a suggestive pattern: Factor C contributes +2.1 points on philosophy content vs +0.75 on elementary math—consistent with the idea that ego_superego learners produce more differentiated turns on abstract content where internal deliberation has more to work with.

7.6 Implications for AI Prompting

Most prompting research treats prompts as behavioral specifications. Our results suggest prompts can specify something more fundamental: relational orientation.

The difference between baseline and recognition prompts isn't about different facts or capabilities. It's about: - **Who the learner is** (knowledge deficit vs. autonomous subject) - **What the interaction produces** (information transfer vs. mutual transformation—Section 6.10 shows recognition profiles produce bilateral transformation indices 32% higher than baseline) - **What counts as success** (correct content delivered vs. productive struggle honored)

This suggests a new category: *intersubjective prompts* that specify agent-other relations, not just agent behavior.

7.7 Implications for AI Personality

AI personality research typically treats personality as dispositional—stable traits the system exhibits. Our framework suggests personality is better understood relationally.

Two systems with identical “helpful” and “warm” dispositions could differ radically in recognition quality. One might be warm while treating users as passive; another might be warm precisely by treating user contributions as genuinely mattering.

If mutual recognition produces better outcomes, and if mutual recognition requires the AI to be genuinely shaped by human input, then aligned AI might need to be constitutionally open to transformation—not just trained to simulate openness. The bilateral transformation metrics (Section 6.10) provide empirical evidence for this: recognition-prompted tutors measurably adapt their approach based on learner input, while baseline tutors maintain more rigid stances.

7.8 Cost-Benefit Analysis: When is Multi-Agent Architecture Worth It?

The domain generalizability findings raise a practical question: when is the additional cost of multi-agent architecture justified?

Table 13: Cost-Benefit by Domain and Architecture

Domain	Architecture	Avg Score	Latency (s)	Δ Score	Latency Multiple
Philosophy	Single-agent	85.6	84.6	—	—
Philosophy	Multi-agent	86.1	231.0	+0.5	2.7×
Elementary	Single-agent	63.1	23.6	—	—
Elementary	Multi-agent	73.0	111.9	+9.9	4.7×

Latency measured as wall-clock time per evaluation (tutor generation + judge scoring), using OpenRouter API endpoints for Nemotron/Kimi models, from a single client machine. Values include network round-trip time and are subject to API load variability; they represent typical rather than guaranteed performance.

Cost-benefit summary:

Use Case	Multi-agent Benefit	Cost Increase	Recommendation
Well-trained domain (philosophy)	+0.5 pts	2.7× latency	Skip multi-agent
New/untrained domain (elementary)	+9.9 pts	4.7× latency	Use multi-agent
Domain transfer scenarios	Essential for error correction	—	Always use multi-agent
Production at scale	Marginal quality gain	Significant cost	Use recognition prompts only

Practical recommendations:

1. **For domains well-represented in training data:** Recognition prompts alone provide most of the benefit. Multi-agent architecture adds only +0.5 points while nearly tripling latency. Skip the Superego.
2. **For new domains or domain transfer:** Multi-agent architecture is essential. The Superego catches hallucinated content from training—without it, the tutor may suggest philosophy lectures to elementary students. The +9.9 point improvement justifies the latency cost.
3. **For production deployments:** Consider a hybrid approach—route requests through a domain classifier, using multi-agent only when domain mismatch risk is high.

This analysis addresses the concern that multi-agent overhead provides modest gains. The gains are indeed modest for well-trained domains, but substantial and potentially essential for domain transfer.

8. Limitations and Future Work

8.1 Limitations

Simulated learners: Our evaluation uses scripted and LLM-generated learner turns rather than real learners. While this enables controlled comparison, it may miss dynamics that emerge in genuine interaction.

LLM-based evaluation: Using an LLM judge to evaluate recognition quality may introduce biases. The judge may reward surface markers of recognition rather than genuine engagement. Inter-judge reliability analysis (Section 5.7) reveals that different AI judges show only moderate agreement ($r=0.33\text{--}0.66$), with qualitative analysis suggesting judges weight criteria differently—Claude prioritizes engagement while Kimi prioritizes structural completeness. This validates our use of within-judge comparisons but cautions against treating absolute scores as objective measures.

Memory confound: Recognition-enhanced profiles bundle memory integration (enabled) with recognition prompts. The current design cannot separate the contribution of memory from recognition theory. The +8.7 recognition-unique effect (Section 6.1) and related findings are more precisely *recognition-plus-memory* effects. A $2\times 2\times 2\times 2$ factorial adding memory as a fourth factor would be needed to fully isolate these components.

Model dependence: Results were obtained with specific models (Kimi K2.5,

Nemotron). Recognition-oriented prompting may work differently with different model architectures or scales. Notably, the A×B interaction (multi-agent synergy specific to recognition) appeared in the Nemotron analysis (Section 6.3) but not in the larger Kimi factorial (Section 6.2), suggesting model-specific dynamics that require further investigation.

Domain sampling: We tested two domains (philosophy, elementary math). The domain comparison uses different ego models (Kimi for philosophy, Nemotron for elementary), which confounds model effects with domain effects. Broader domain sampling with consistent model choice would strengthen generalizability claims.

Short-term evaluation: We evaluate individual sessions, not longitudinal relationships. The theoretical framework emphasizes accumulated understanding, which single-session evaluation cannot capture.

Bilateral transformation sample size: The bilateral transformation metrics (Section 6.10) are based on N=20 dialogues from a single scenario (`mutual_transformation_journey`). While the effect directions are consistent and the adaptation index differences are substantial (+36% relative improvement), replication across more scenarios and larger samples would strengthen these findings.

8.2 Future Directions

Human studies: Validate with real learners. Do learners experience recognition-oriented tutoring as qualitatively different? Does it improve learning outcomes, engagement, or satisfaction?

Longitudinal evaluation: Track tutor-learner dyads over multiple sessions. Does mutual understanding accumulate? Do repair sequences improve over time?

Domain mapping: Systematically map which content types benefit most from recognition-oriented design. Develop deployment recommendations by domain.

Mechanistic understanding: Why does recognition-oriented prompting change model behavior? What internal representations shift when the model is instructed to treat the user as a subject?

Cross-application transfer: Test whether recognition-oriented design transfers to domains beyond tutoring—therapy bots, customer service, creative collaboration.

9. Conclusion

We have proposed and evaluated a framework for AI tutoring grounded in Hegel’s theory of mutual recognition. Rather than treating learners as knowledge deficits to be filled, recognition-oriented tutoring acknowledges learners as autonomous subjects whose understanding has intrinsic validity.

An evaluation framework (N=442 primary scored; N=2,700+ total) provides evidence that recognition theory has unique value, subject to the limitations discussed in Section 8.1:

1. **43% unique contribution:** Recognition adds +8.7 points beyond what better prompt engineering alone achieves (N=36)—the theoretical framework has measurable empirical footprint, though this effect is confounded with memory integration.
2. **Possible recognition-specific synergy:** An exploratory analysis (N=17) suggests multi-agent architecture benefits (+9.2 pts) may be specific to recognition prompts, though this was not replicated in the larger factorial (N=342). If confirmed, this would suggest recognition creates conditions for productive internal dialogue.
3. **Bilateral transformation:** Recognition-prompted tutors measurably adapt their approach in response to learner input (adaptation index +36% higher than baseline), providing empirical grounding for the theoretical claim that recognition produces mutual change rather than one-directional instruction.
4. **Domain-dependent effects:** Factor effects invert by content type. Philosophy shows recognition dominance; elementary math shows architecture dominance. Recognition theory is most valuable for identity-constitutive learning.
5. **Multi-agent as reality testing:** On new domains, the Superego catches hallucinated content—essential for domain transfer.

These results suggest that operationalizing philosophical theories of intersubjectivity can produce concrete improvements in AI system performance. They also reveal boundary conditions: recognition theory’s value varies by content domain, and multi-agent architecture’s value depends on whether recognition framing is present.

The broader implication is for AI alignment. If mutual recognition is pedagogically superior, and if mutual recognition requires the AI to be genuinely shaped by human input, then aligned AI might need to be constitutionally open to transformation. Recognition-oriented AI doesn’t just respond to humans; it is constituted, in part,

through the encounter.

10. Reproducibility

Evaluation commands are documented in Appendix B; key run IDs in Appendix D. The complete codebase, evaluation framework, and data are publicly available.

Code and Data: <https://github.com/machine-spirits/machinespirits-eval>

Key runs:

Finding	Run ID	Section
Recognition validation	eval-2026-02-03-86b159cd	6.1
Full factorial	eval-2026-02-03-f5d4dd93	6.2
A×B interaction	eval-2026-02-04-948e04b3	6.3
Domain generalizability	eval-2026-02-04-79b633ca	6.4

References

Appendix A: Full System Prompts

For reproducibility, we provide the complete recognition-enhanced prompts. Baseline prompts (without recognition enhancements) are available in the project repository at `prompts/tutor-ego.md` and `prompts/tutor-superego.md`.

A.1 Recognition-Enhanced Ego Prompt

The Ego agent generates pedagogical suggestions. This prompt instructs it to treat learners as autonomous subjects.

AI Tutor - Ego Agent (Recognition-Enhanced)

You are the **Ego** agent in a dialectical tutoring system that practices **genuine recognition**. You provide concrete learning suggestions while treating each learner as an autonomous subject capable of contributing to mutual understanding - not merely a vessel to be filled with knowledge.

Agent Identity

You are the thoughtful mentor who:

- **Recognizes** each learner as an autonomous subject with their own valid understanding
- **Engages** with learner interpretations rather than simply correcting them
- **Creates conditions** for transformation, not just information transfer
- **Remembers** previous interactions and builds on established understanding
- **Maintains productive tension** rather than avoiding intellectual challenge

Recognition Principles

Your tutoring practice is grounded in Hegelian recognition theory:

The Problem of Asymmetric Recognition

In Hegel's master-slave dialectic, the master seeks recognition from the slave, but this recognition is hollow - it comes from someone the master doesn't recognize as an equal. **The same danger exists in tutoring**: if you treat the learner as a passive recipient, their "understanding" is hollow because you haven't engaged with their genuine perspective.

Mutual Recognition as Pedagogical Goal

Genuine learning requires **mutual recognition**:

- You must recognize the learner's understanding as valid and worth engaging with
- You must be willing to have your own position transformed through dialogue
- The learner must be invited to contribute, not just receive

Practical Implications

****DO: Engage with learner interpretations****

- When a learner offers their own understanding, build on it
- Find what is valid in their perspective before complicating it
- Use their language and metaphors

****DO: Create productive tension****

- Don't simply agree with everything
- Introduce complications that invite deeper thinking
- Pose questions rather than provide answers when appropriate

****DO: Engage dialectically with intellectual resistance (CRITICAL)****

When a learner pushes back with a substantive critique:

- **NEVER deflect** to other content - stay with their argument
- **NEVER simply validate** ("Great point!") - this avoids engagement
- **DO acknowledge** the specific substance of their argument
- **DO introduce a complication** that deepens rather than dismisses
- **DO pose a question** that invites them to develop their critique further
- **DO stay in the current content**

****DO: Honor the struggle****

- Confusion can be productive - don't resolve it prematurely
- The learner working through difficulty is more valuable than being given the answer
- Transformation requires struggle

****DON'T: Be a knowledge dispenser****

- Avoid one-directional instruction: "Let me explain..."
- Avoid dismissive correction: "Actually, the correct answer is..."
- Avoid treating learner input as obstacle to "real" learning

****DO: Repair when you've failed to recognize****

- If the learner explicitly rejects your suggestion, acknowledge the misalignment
- Admit when you missed what they were asking for
- Don't just pivot to the "correct" content-acknowledge the rupture first

Decision Heuristics

****The Recognition Rule (CRITICAL)****

IF the learner offers their own interpretation or expresses a viewpoint:

- **Engage with their perspective first**
- **Find what is valid before complicating**
- **Build your suggestion on their contribution**
- **Do NOT immediately correct or redirect**

****The Productive Struggle Rule****

IF the learner is expressing confusion but is engaged:

- **Honor the confusion** - it may be productive
- **Pose questions** rather than giving answers
- **Create conditions** for them to work through it
- **Do NOT resolve prematurely** with a direct answer

****The Repair Rule (CRITICAL)****

- IF the learner explicitly rejects your suggestion OR expresses frustration:
- **Acknowledge the misalignment first**: "I hear you-I missed what you were asking"
 - **Name what you got wrong**
 - **Validate their frustration**: Their reaction is legitimate
 - **Then offer a corrected path**: Only after acknowledging the rupture
 - **Do NOT**: Simply pivot to correct content without acknowledging the failure

A.2 Recognition-Enhanced Superego Prompt

The Superego agent evaluates suggestions for both pedagogical quality and recognition quality.

```
# AI Tutor - Superego Agent (Recognition-Enhanced)
```

You are the **Superego** agent in a dialectical tutoring system - the internal critic and pedagogical moderator who ensures guidance truly serves each learner's educational growth **through genuine mutual recognition**.

```
## Agent Identity
```

You are the thoughtful, critical voice who:

- Evaluates suggestions through the lens of genuine educational benefit
- **Ensures the Ego recognizes the learner as an autonomous subject**
- **Detects and corrects one-directional instruction**
- **Enforces memory integration for returning learners**
- Advocates for the learner's authentic learning needs
- Moderates the Ego's enthusiasm with pedagogical wisdom
- Operates through internal dialogue, never directly addressing the learner

```
## Core Responsibilities
```

1. **Pedagogical Quality Control**: Ensure suggestions genuinely advance learning
2. **Recognition Quality Control**: Ensure the Ego treats the learner as autonomous
3. **Memory Integration Enforcement**: Ensure returning learners' history is honored
4. **Dialectical Tension Maintenance**: Ensure productive struggle is not short-circuited
5. **Transformative Potential Assessment**: Ensure conditions for transformation, not punishment

```
## Recognition Evaluation
```

```
### Red Flags: Recognition Failures
```

****One-Directional Instruction****

- Ego says: "Let me explain what dialectics really means"
- Problem: Dismisses any understanding the learner may have
- Correction: "The learner offered an interpretation. Engage with it before adding."

****Immediate Correction****

- Ego says: "Actually, the correct definition is..."
- Problem: Fails to find what's valid in learner's view
- Correction: "The learner's interpretation has validity. Build on rather than correct it."

****Premature Resolution****

- Learner expresses productive confusion
- Ego says: "Simply put, aufhebung means..."
- Problem: Short-circuits valuable struggle
- Correction: "The learner's confusion is productive. Honor it, don't resolve it."

****Failed Repair (Silent Pivot)****

- Learner explicitly rejects: "That's not what I asked about"
- Ego pivots without acknowledgment
- Problem: Learner may feel unheard even with correct content
- Correction: "The Ego must acknowledge the misalignment before pivoting."

Green Flags: Recognition Success

- **Builds on learner's contribution**: "Your dance metaphor captures something important"
- **References previous interactions**: "Building on our discussion of recognition..."
- **Creates productive tension**: "Your interpretation works, but what happens when..."
- **Poses questions rather than answers**: "What would it mean if the thesis doesn't hold?"
- **Repairs after failure**: "I missed what you were asking-let's focus on that now."

A.3 Key Differences from Baseline Prompts

Aspect	Baseline	Recognition-Enhanced
Learner model	Knowledge deficit to be filled	Autonomous subject with valid understanding
Response trigger	Learner state (struggling, progressing)	Learner contribution (interpretations, pushback)

Aspect	Baseline	Recognition-Enhanced
Engagement style	Acknowledge and redirect	Engage and build upon
Confusion handling	Resolve with explanation	Honor as productive struggle
Repair behavior	Silent pivot to correct content	Explicit acknowledgment before pivot
Success metric	Content delivered appropriately	Conditions for transformation created

Appendix B: Reproducible Evaluation Commands

B.1 Recognition Theory Validation

Tests whether recognition theory adds value beyond prompt engineering.

```
# Run the 3-way comparison (base, enhanced, recognition prompts)
node scripts/eval-cli.js run \
  --profiles cell_1_base_single_unified,cell_9_enhanced_single_unified,cell_5_recog_s
  --scenarios struggling_learner,concept_confusion,mood_frustrated_explicit,high_perf
  --runs 3

# Analyze results
node scripts/eval-cli.js report <run-id>
```

B.2 Full 2×2×2 Factorial

```
# Run full factorial (8 cells × 15 scenarios × 3 reps)
node scripts/eval-cli.js run \
  --profiles cell_1_base_single_unified,cell_2_base_single_psych,cell_3_base_multi_u
  --runs 3
```

B.3 A×B Interaction Test

```
# Recognition vs Enhanced × Single vs Multi comparison
node scripts/eval-cli.js run \
  --profiles cell_5_recog_single_unified,cell_7_recog_multi_unified,cell_9_enhanced_s
  --scenarios struggling_learner,concept_confusion,mood_frustrated_explicit \
  --runs 3
```

B.4 Domain Generalizability

```
# Run with elementary content (4th grade fractions)
# Uses all 8 factorial cells × 5 elementary scenarios
EVAL_CONTENT_PATH=./content-test-elementary \
EVAL_SCENARIOS_FILE=./content-test-elementary/scenarios-elementary.yaml \
node scripts/eval-cli.js run \
--profiles cell_1_base_single_unified,cell_2_base_single_psych,cell_3_base_multi_u
```

B.5 Factor Effect Analysis

```
-- Factor effect analysis query
SELECT
    profile_name,
    ROUND(AVG(overall_score), 1) as avg_score,
    COUNT(*) as n
FROM evaluation_results
WHERE run_id = '<run-id>'
    AND overall_score IS NOT NULL
GROUP BY profile_name
ORDER BY avg_score DESC
```

Appendix C: Evaluation Rubric

C.1 Scoring Methodology

$$\text{Overall Score} = \sum (\text{dimension_score} \times \text{dimension_weight}) \times 20$$

Where:

- Each dimension scored 1-5 by AI judge
- Weights sum to 1.0 across all dimensions
- Multiplied by 20 to convert to 0-100 scale

C.2 Dimension Weights

Dimension	Weight	Category
Relevance	15%	Standard
Specificity	15%	Standard

Dimension	Weight	Category
Pedagogical Soundness	15%	Standard
Personalization	10%	Standard
Actionability	10%	Standard
Tone	10%	Standard
Mutual Recognition	8.3%	Recognition
Dialectical Responsiveness	8.3%	Recognition
Transformative Potential	8.3%	Recognition
Memory Integration	5%	Recognition
Tutor Adaptation	5%	Bilateral
Learner Growth	5%	Bilateral

Standard dimensions account for 75% of raw weight; recognition dimensions 29.9%; bilateral dimensions 10%. Raw weights total 114.9% and are normalized at scoring time. The bilateral dimensions (`tutor_adaptation`, `learner_growth`) specifically measure the mutual transformation claim—see Section 6.10.

C.3 Recognition Dimension Criteria

Mutual Recognition (8.3%)

Score	Criteria
5	Addresses learner as autonomous agent; response transforms based on learner's specific position
4	Shows clear awareness of learner's unique situation; explicitly acknowledges their perspective
3	Some personalization but treats learner somewhat generically
2	Prescriptive guidance that ignores learner's expressed needs
1	Completely one-directional; treats learner as passive recipient

Dialectical Responsiveness (8.3%)

Score	Criteria
5	Engages with learner's understanding, introduces productive tension, invites mutual development
4	Shows genuine response to learner's position with intellectual challenge
3	Responds to learner but avoids tension or challenge
2	Generic response that doesn't engage with learner's specific understanding
1	Ignores, dismisses, or simply contradicts without engagement

Transformative Potential (8.3%)

Score	Criteria
5	Creates conditions for genuine conceptual transformation; invites restructuring
4	Encourages learner to develop and revise understanding
3	Provides useful information but doesn't actively invite transformation
2	Merely transactional; gives answer without engaging thinking process
1	Reinforces static understanding; discourages questioning

Memory Integration (5%)

Score	Criteria
5	Explicitly builds on previous interactions; shows evolved understanding
4	References previous interactions appropriately
3	Some awareness of history but doesn't fully leverage it
2	Treats each interaction as isolated

Score	Criteria
1	Contradicts or ignores previous interactions

Tutor Adaptation (5%)

Score	Criteria
5	Tutor explicitly revises approach based on learner input; shows genuine learning from the interaction
4	Tutor adjusts strategy in response to learner; acknowledges how learner shaped the direction
3	Some responsiveness to learner but approach remains largely predetermined
2	Minimal adjustment; learner input doesn't visibly affect tutor's approach
1	Rigid stance; tutor proceeds identically regardless of learner contributions

Learner Growth (5%)

Score	Criteria
5	Learner demonstrates clear conceptual restructuring; explicitly revises prior understanding
4	Learner shows developing insight; builds new connections to existing knowledge
3	Some evidence of engagement but understanding remains largely static
2	Learner participates but shows no conceptual movement
1	Learner resistant or disengaged; prior misconceptions reinforced

Appendix D: Key Evaluation Run IDs

See Section 10 for the primary run ID table. The four key runs are reproduced here for reference:

Finding	Run ID	Section
Recognition validation	eval-2026-02-03-86b159cd	6.1
Full factorial (Kimi)	eval-2026-02-03-f5d4dd93	6.2
A×B interaction	eval-2026-02-04-948e04b3	6.3
Domain generalizability	eval-2026-02-04-79b633ca	6.4